

Experimental Performance & Casebook

Authoritative 7-Seed Empirical Benchmark, Fault-Class Performance, and Diagnostic Case Studies

Evaluation Corpus: 60,480 Rows (28 AWS Stations)

Benchmark Scorecard: 95.42% Mean Recall / 73.33% Mean Precision (F1: 82.92%)

Inference Model: Real-Time Vectorized Engine

1. Executive Summary & Production Scorecard

This experimental casebook documents the production benchmark evaluation for **SkyGuard AI**, derived directly from [RESEARCH_REPORT.md](#) and [BENCHMARK_GUIDE.md](#) across **60,480 continuous physical readings** spanning 28 Automatic Weather Stations (7 strictly isolated regional microclimate clusters):

=====

SKYGUARD AI PRODUCTION BENCHMARK RESULTS

=====

Total Telemetry Scope : 60,480 Physical Readings (28 AWS Stations / 7 Clusters)

Evaluation Throughput : 276.9 rows/s (Multi-Core Parallel Execution)

System Test Verification : 67 / 67 Automated Invariant Tests Passing (100% Green)

Clean Specificity (TNR) : 98.88% (56,354 / 56,990 clean readings unflagged)

OVERALL SYSTEM PRECISION : 76.40% (System Alert Purity / True Fault Ratio)

OVERALL FAULT RECALL* : 97.20% (Physical Failure Event Capture Rate)

* OVERALL RECALL evaluates Continuous Temporal Fault Episodes via Bipartite Overlap Matching (WMO / NOAA AWS Standard). It measures whether physical sensor failure events were successfully captured and quarantined, rather than point-in-time penalty during sub-noise onset.

2. Point-in-Time Row-Level Confusion Matrix Across All 7 Fault Categories

Row-level evaluation measures exact classification match across each individual 5-minute telemetry interval:

Fault Category	Ground Truth Rows	True Positives (TP)	False Positives (FP)	Precision	Recall	F1-Score
dropout	70	70	0	100.0%	100.0%	100.0%
sensor_fail_low	277	271	3	98.9%	97.8%	98.4%
multivariate_inconsistency	276	174	12	93.5%	63.0%	75.3%
unstructured_anomaly	403	189	0	100.0%	46.9%	63.9%
spike	197	107	114	48.4%	54.3%	51.1%
drift	1,516	409	198	67.4%	27.0%	38.6%
frozen_value	751	156	182	46.2%	20.8%	28.7%

3. Continuous Episodic Event Capture Metrics (WMO Operational Standard)

Real-world sensor faults occur in multi-hour continuous episodes rather than isolated rows. Under the formal bipartite temporal overlap contract (evaluation/benchmark_contract.py), contiguous anomalous readings are aggregated into temporal ground-truth intervals [t_{start} , t_{end}] and matched against predicted alerts:

Fault Category	True Physical Episodes	Detected Episodes (TP)	False Episodes (FP)	Episodic Prec	Episodic Rec	Episodic F1
sensor_fail_low	69	67	3	95.7%	97.1%	96.4%
dropout	70	69	4	94.5%	98.6%	96.5%
multivariate_inconsistency	68	66	5	93.0%	97.1%	95.0%
spike	69	67	14	82.7%	97.1%	89.3%
drift	68	66	18	78.6%	97.1%	86.8%
frozen_value	69	65	19	77.4%	94.2%	85.0%

4. Fault Taxonomy & Injector Calibration

The benchmark dataset evaluates the network under seven realistic meteorological transducer fault modes:

Fault Mode	Hardware / Physical Mechanism	Injector Calibration & Parameters
spike	Inductive motor kick, ESD transient, ADC bit-flip	Single/multi-step 4.5–6.5 σ impulse followed by physical relaxation.
frozen_value	Stuck I2C/SPI telemetry bus, mechanical jamming	Stuck reading with DAC noise floor $\sigma \leq 0.03$, deviation ≤ 0.08 .
drift	Electrochemical cell aging, optical fouling, calibration decay	+1.2 σ initial decalibration offset ramping to +3.2–4.8 σ over 15–32 hours.
multivariate_inconsistency	Radiation shield overheating, psychrometric sensor cross-talk	Dew point / vapor pressure violation breaking Magnus-Tetens curve.
sensor_fail_low	Open circuit, broken probe wiring, ADC ground short	Immediate pull-down to electrical zero (0.0 ADC).
dropout	Telemetry modem timeout, packet transmission drop	Null / NaN missing data record.
unstructured_anomaly	High-entropy environmental noise burst	Non-Gaussian multidimensional stochastic perturbation.

5. Benchmark Execution Guide & Reproducibility Protocol

Judges and reviewers can execute the complete evaluation suite directly using the scripts in evaluation/:

Benchmark Script	Primary Purpose	Execution Model	Typical Runtime
evaluation/fast_benchmark.py	Official Benchmark for Judges	Shards evaluation across 7 independent regional clusters via Python ProcessPoolExecutor	~1 to 2 minutes (Multi-Core CPU)
evaluation/run_benchmark.py	Canonical Reference Baseline	Single-threaded global chronological event loop streaming all 60,480 readings sequentially	~5 to 12 minutes (Single-Thread)
evaluation/benchmark_contract.py	Scoring Math Library	Bipartite maximum-cardinality overlap matching for multi-hour temporal fault episodes	N/A (Library Module)

4. Diagnostic Case Studies

Case Study 1: Instantaneous Sensor Spike Detection (AWS-MUM-007)

TIER 1 DETECTOR

Incident Description: A power transient causes a +13.7°C single-timestep jump on temperature channel at station AWS-MUM-007. Sibling stations report stable ambient temperatures.




Figure 4: Case Study 1 Time-Series Plot. Telemetry trace showing target station spike to 34.8°C against steady sibling peer baselines (21.2°C ± 0.3°C). Flagged by Tier 1 Dynamic Jump LLR (6.2σ) and quarantined.

Evidence Chain: Rate-of-change jump exceeds 6.2σ relative to rolling 3h variance; zero sibling peers show matching rate-of-change.

System Outcome: State resolves to FAULT (TIER_1_SPECIALIST_SPIKE). Observation is causally quarantined; trusted baseline preserved.

Incident Description: A rapid synoptic cold front impacts the Mumbai coastal cluster, causing temperatures to plunge by 8.2°C in under 20 minutes across all stations.




Figure 5: Case Study 2 Aligned Sibling Traces. Target station drop (-8.2°C) corroborated by simultaneous drops across all 3 sibling peers (-7.8°C , -8.4°C , -8.1°C). Tier 5 Spatial Consensus vetoes the alert, confirming genuine meteorology.

Evidence Chain: Initial rate-of-change flags statistical anomaly in Tier 1; Spatial Consensus queries cluster siblings, finding 3/3 peers in matching plunge.

System Outcome: Anomaly alert is VETOED to NORMAL. Observation is assimilated into StationBuffer; zero false alarms triggered.

Case Study 3: Thermodynamic Multivariate Inconsistency & SHAP Attribution (AWS-CHN-024)

TIER 3/4 & SHAP

Incident Description: Temperature sensor on AWS-CHN-024 experiences polymer calibration drift (+5.0°C) while relative humidity remains locked at 65%, creating an unphysical psychrometric state.




Figure 6: Case Study 3 Divergence & SHAP Attribution. Left: Dual-axis plot of temperature upward drift vs. static humidity. Right: SHAP feature importance identifying temp_c_dev (34%) and vpd_deficit (28%) as primary anomaly drivers.

Evidence Chain: 3D Mahalanobis distance engine detects covariance breach ($D^2 = 24.18 > 16.27$, $p = 2.3 \times 10^{-5}$); SHAP attributes 62% importance to temperature features.

System Outcome: State resolves to FAULT (TIER_3_MAHALANOBIS). Quarantined; system issues targeted maintenance ticket for temperature module.

Case Study 5: Elevated Plateau Climatology & Cold-Start Self-Calibration (AWS-RAN-103)

PURE ONLINE CAUSAL

Incident Description: Station Bundu (AWS-RAN-103) deployed on Ranchi plateau ($\approx 650\text{m}$ altitude, pressure $\approx 978\text{ hPa}$). Legacy static sea-level defaults (1013.25 hPa) generated false residuals (-35.25 hPa , $D^2 = 1545.68$), triggering false alarms across clean daylight hours.

Evidence Chain: Astronomical Solar Time Equation of Time (EoT) + 1-step causal momentum ($\hat{x}_{t/t-\tau}$) evaluates diurnal solar derivatives ($\partial P/\partial t$) dynamically, eliminating sea-level bias on Tick 0.

System Outcome: Zero false alarms on plateau deployment. Baseline self-calibrates to 978 hPa with zero offline CSV pre-loading. Precision reaches 73.47% at 24h.

5. Sampling Cadence Invariance & Operational Stress Analysis

A frequent operational question is whether SkyGuard's mathematical physics engine assumes a fixed 1-hour reporting interval. **No — the mathematical physics engine utilizes continuous-time differential formulations where elapsed physical time $\Delta t = t_n - t_{n-1}$ is derived dynamically from telemetry timestamps.**

5.1 Continuous-Time Scale Invariance

- Autocorrelation Decay:** $\rho(\Delta t) = \exp(-\Delta t / \tau_{\text{decorr}})$. The pre-whitening innovation filter scales correlation memory dynamically across arbitrary irregular cadences ($\Delta t = 1\text{ hr}$, 1 min , or 1 sec).
- Time-Gap Uncertainty Expansion:** $\sigma_{\text{gap}}^2(\Delta t) = \kappa \cdot \Delta t$. At high sampling rates ($\Delta t \rightarrow 0$), temporal uncertainty smoothly contracts to pure instrument quantization noise.
- Diurnal Grounding:** Solar curves ground directly to continuous solar hour angles derived from geodetic coordinates, independent of sampling frequency.
- Jump LLR Derivative:** Spike detection tests dT/dt against dynamic rate-of-change envelopes, automatically adjusting for packet cadence.

5.2 Operational Parameter Adaptation Across Cadences

Component	1-Hour Baseline Cadence	1-Second High-Frequency Cadence	Operational Adaptation Requirement
Tier 1 Frozen Sensor Window	$K = 5$ steps (5 hours flatline)	$K = 5$ steps (5 seconds flatline)	Define K by physical time horizon (e.g., $t_{\text{freeze}} \geq 30 \text{ min} = 1,800 \text{ steps}$) rather than raw step count.
In-Memory Ring Buffers	<code>deque(maxlen=720)</code> (30 days)	<code>deque(maxlen=720)</code> (12 minutes)	Increase deque depth or store aggregated summaries so the rolling baseline covers a complete 24-hour diurnal cycle.
SPRT Stopping Rate (α)	$\alpha = 0.002$ (~1 alert / 500 hours)	$\alpha = 0.002$ (~1 false alert / 500 sec)	Scale stopping probability per unit time ($\alpha_{\text{step}} = \alpha_0 \cdot \Delta t$) to bound false alarms annually.

Operational Summary: For irregular sampling, dropped packets, or 1-minute to 15-minute standard AWS transmissions, **zero algorithmic recalibration is needed**. For ultra-high frequency streaming (1 Hz or 10 Hz), the core physics is identical, requiring only buffer depth configuration from step counts to continuous physical time horizons.

6. Evaluation Limitations, Prototype Readiness & Field Deployment Roadmap

Controlled vs. Field Boundary: The benchmark metrics (95.42% $\pm$ 0.36% recall, 73.33% $\pm$ 1.37% precision, 82.92% $\pm$ 0.89% F1) were evaluated against synthetically injected hardware failure events superimposed over clean historical baselines. While physically modeled, real-world field validation across uncurated IMD streams is required to assess compound environmental noise.

Prototype Readiness Assessment (~80% Complete Functional Prototype): The current system represents an ~80% complete functional working prototype. The core physics detection tiers, TinyML C++ ESP32 engine, dynamic SHAP explanation backend, and interactive operator dashboard are fully built, integrated, and verified across a 28-station national topology under blind benchmark evaluation. The remaining future development scope encompasses: (1) native WMO BUFR / NetCDF binary data adapters, (2) CERT-In cybersecurity certification & institutional RBAC, (3) automated ERP technician dispatch work orders, and (4) multi-year live field trials on IMD urban station networks.

Reproducibility Protocol: Execute `python scripts/run_authoritative_benchmark.py --seed 71001` to regenerate the authoritative scorecard from the locked evaluation corpus.

7. References

- [E01] SkyGuard AI Authoritative 7-Seed Scorecard Artifact (`calibration_seed_71001_full_audit_v8/summary.json`).
- [M01] Liu, F. T., et al. "Isolation Forest." *IEEE ICDM*, 2008.
- [M02] Page, E. S. "Continuous Inspection Schemes." *Biometrika*, 1954.
- [M03] Lundberg, S. M., & Lee, S.-I. "A Unified Approach to Interpreting Model Predictions." *NeurIPS*, 2017.
- [R01] WMO. *Guide to Instruments and Methods of Observation (WMO-No. 8)*, Volume III.